



# MapReduce原理

刘军 ( liujun@bupt.edu.cn )

北京邮电大学 数据科学中心

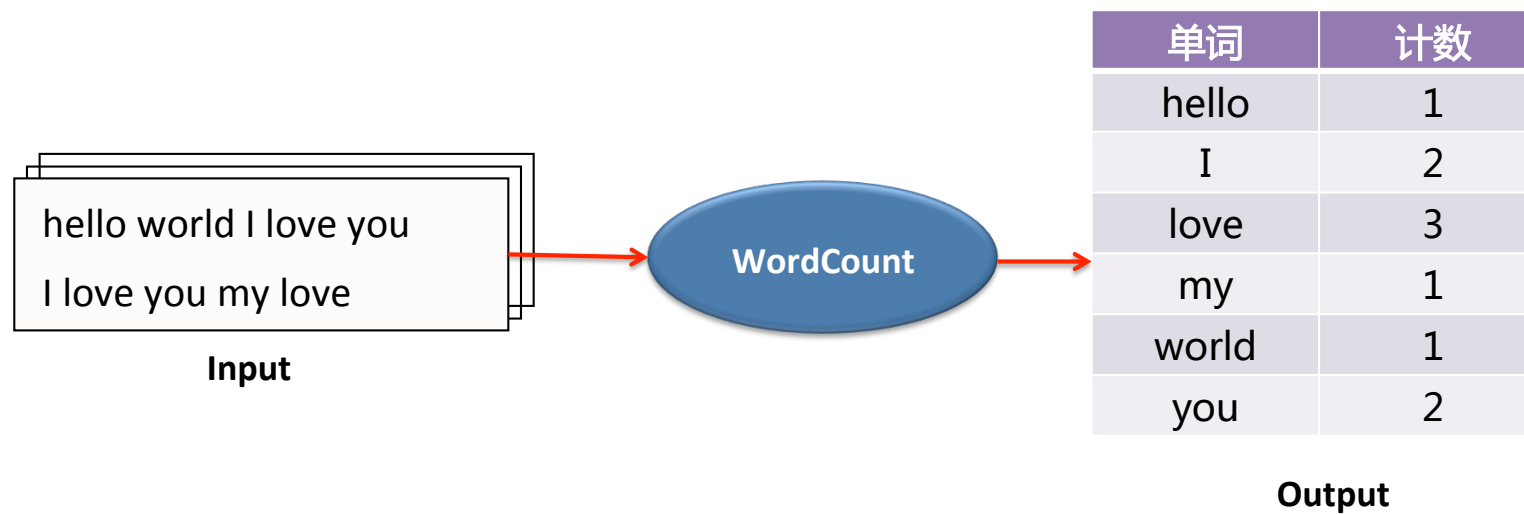


2.0.1稳定版发布

从WordCount开始

## 如何实现以下程序？

---



## 朴素的串行实现

---

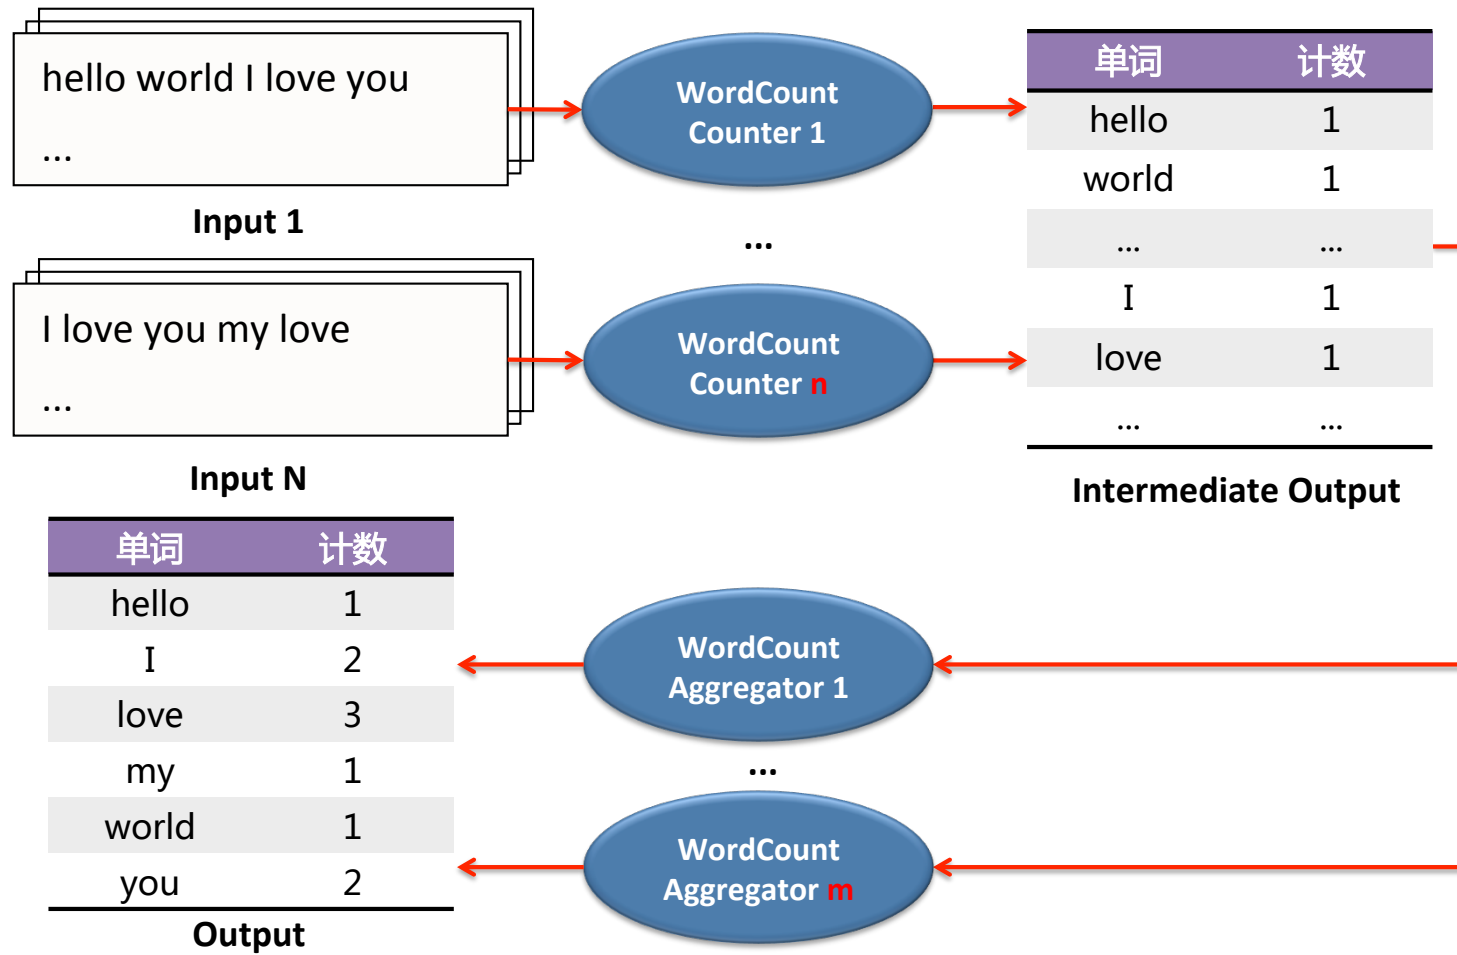
```
1: define wordCount
2:   for each doc in docSet
3:     words = tokenize(doc)
4:     for each word in words
5:       wordCount[word] ++
6:     end for
7:   end for
8:   output(wordCount)
```

- 串行代码的问题：

- 海量数据
- 运行效率

- 怎么办？

## 并行WordCount – 思路



## 并行WordCount – 代码

---

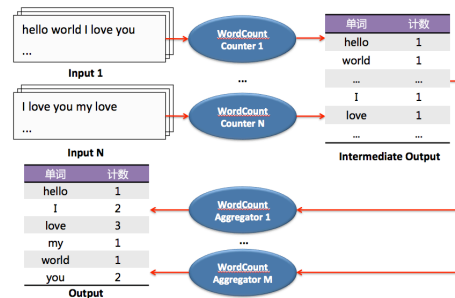
### Counter

```
1: define wordCount
2: for each doc in docSet
3:   words = tokenize(doc)
4:   for each word in words
5:     wordCount[word] ++
6:   end for
7: end for
8: send (wordCount, Aggregator)
```

### Aggregator

```
1: define totalCount
2: for each wordCount from Counter
3:   for each word in wordCount
4:     totalCount[word] += wordCount[word]
5:   end for
6: end for
7: output(totalCount)
```

## 并行WordCount – 架构




**Counter**

```
1: define wordCount
2: for each doc in docSet
3:   words = tokenize(doc)
4:   for each word in words
5:     wordCount[word] ++
6:   end for
7: end for
8: send (wordCount, Aggregator)
```

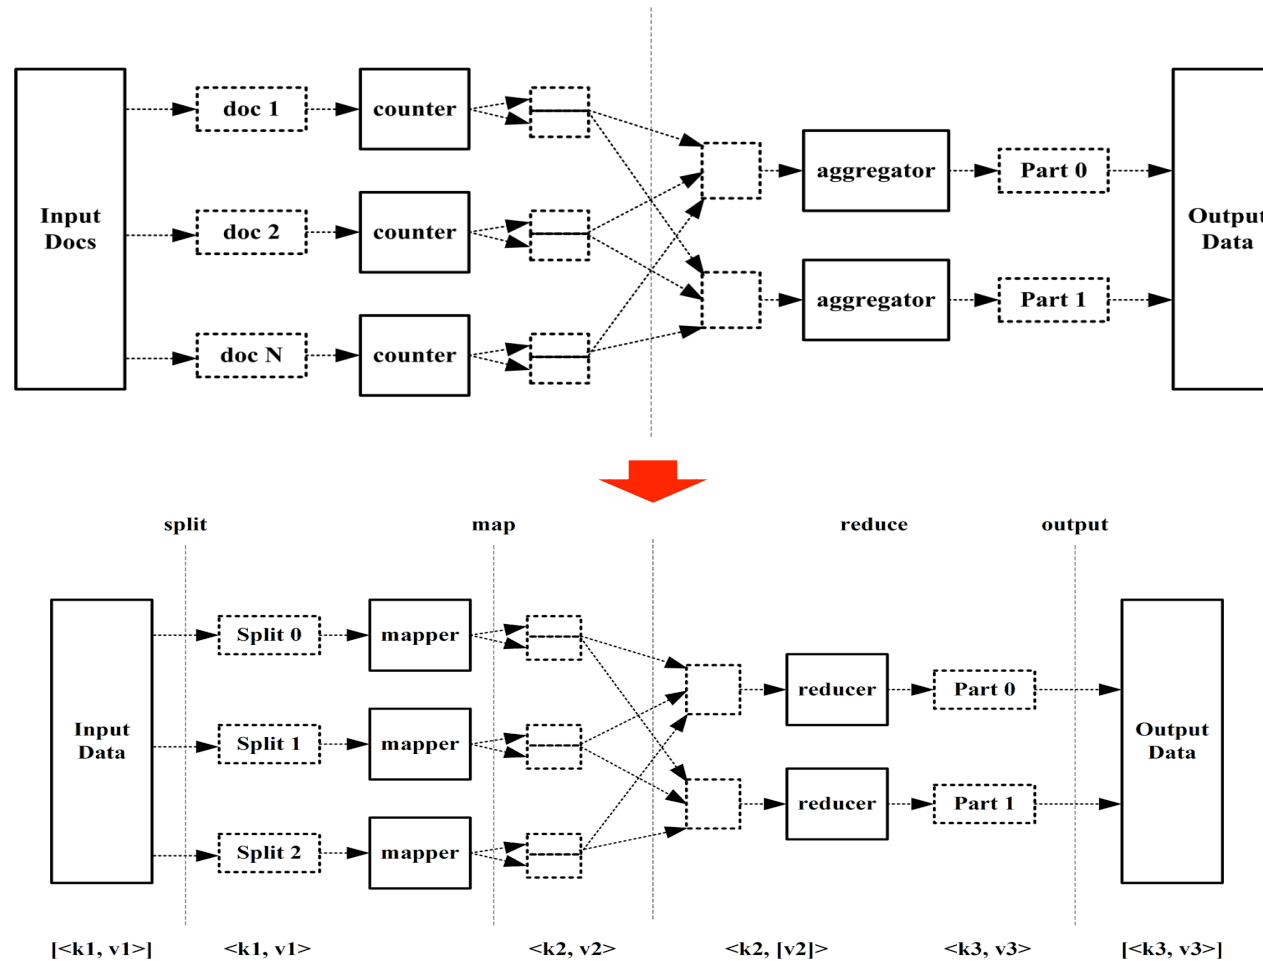
**Aggregator**

```
1: define totalCount
2: for each wordCount from Counter
3:   for each word in wordCount
4:     totalCount[word] += wordCount[word]
5:   end for
6: end for
7: output(totalCount)
```





## 从并行WordCount到MapReduce



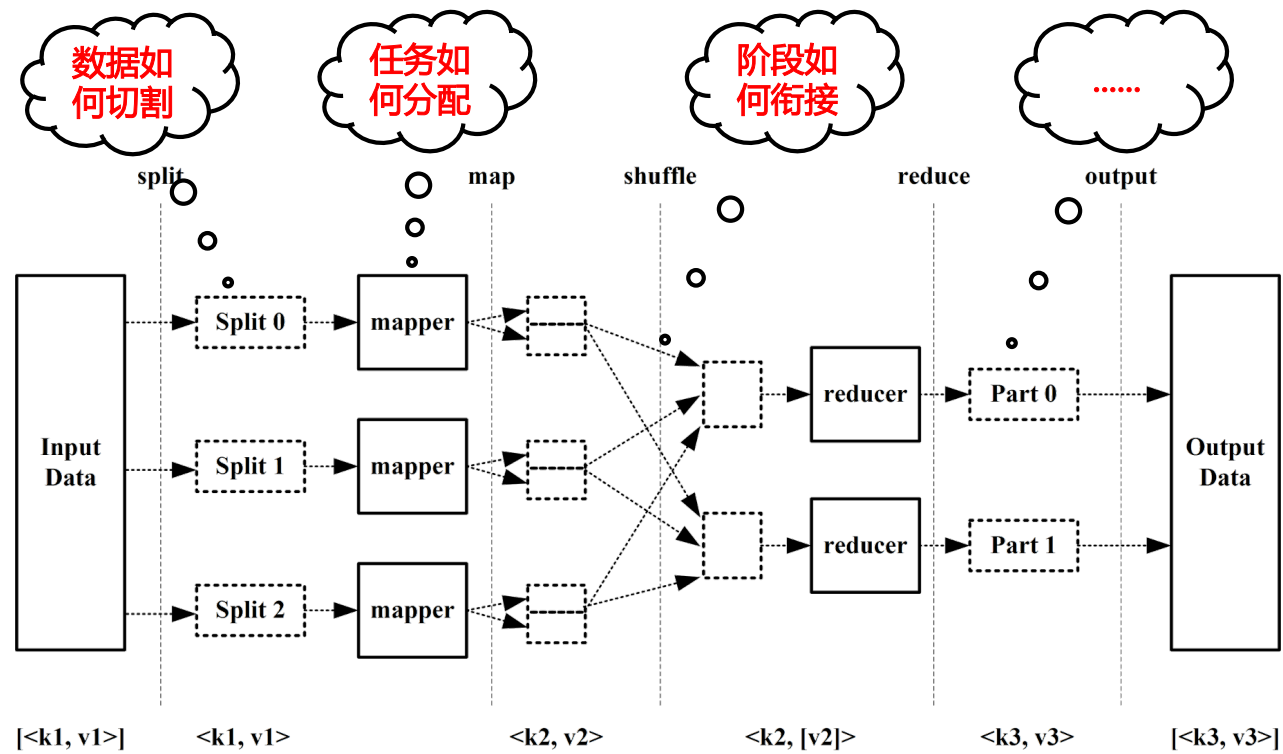
# MapReduce版的WordCount

```
public static class TokenizerMapper extends Mapper
<Object, Text, Text, IntWritable>{
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    public void map(Object key, Text value, Context
        context) throws IOException, InterruptedException{
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

```
public static class IntSumReducer extends Reducer
<Text, IntWritable, Text, IntWritable>{
    private IntWritable result = new IntWritable();
    public void reduce(Text key, Iterable<IntWritable>
        values, Context context) throws IOException{
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

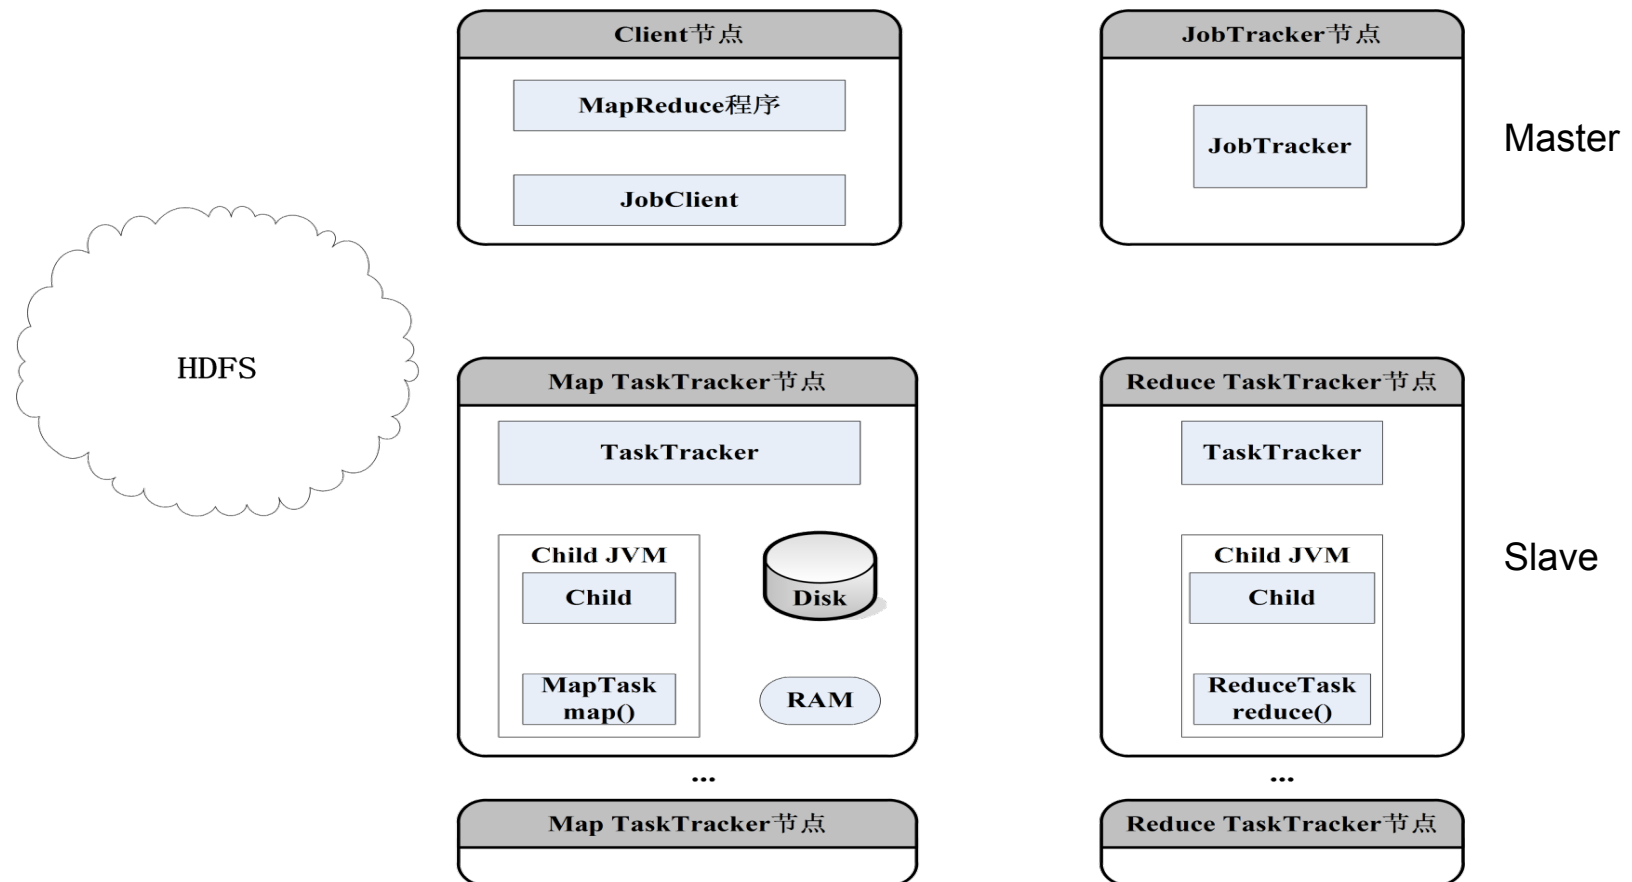
```
public static void main(String[] args) throws Exception{
    Configuration conf = new Configuration();
    String[] otherArgs = new GenericOptionsParser(
        conf, args).getRemainingArgs();
    if (otherArgs.length != 2) {
        System.err.println("Usage: wordcount <in> <out>");
        System.exit(2);
    }
    Job job = new Job(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job,
        new Path(otherArgs[0]));
    FileOutputFormat.setOutputPath(job,
        new Path(otherArgs[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

## 要解决的问题

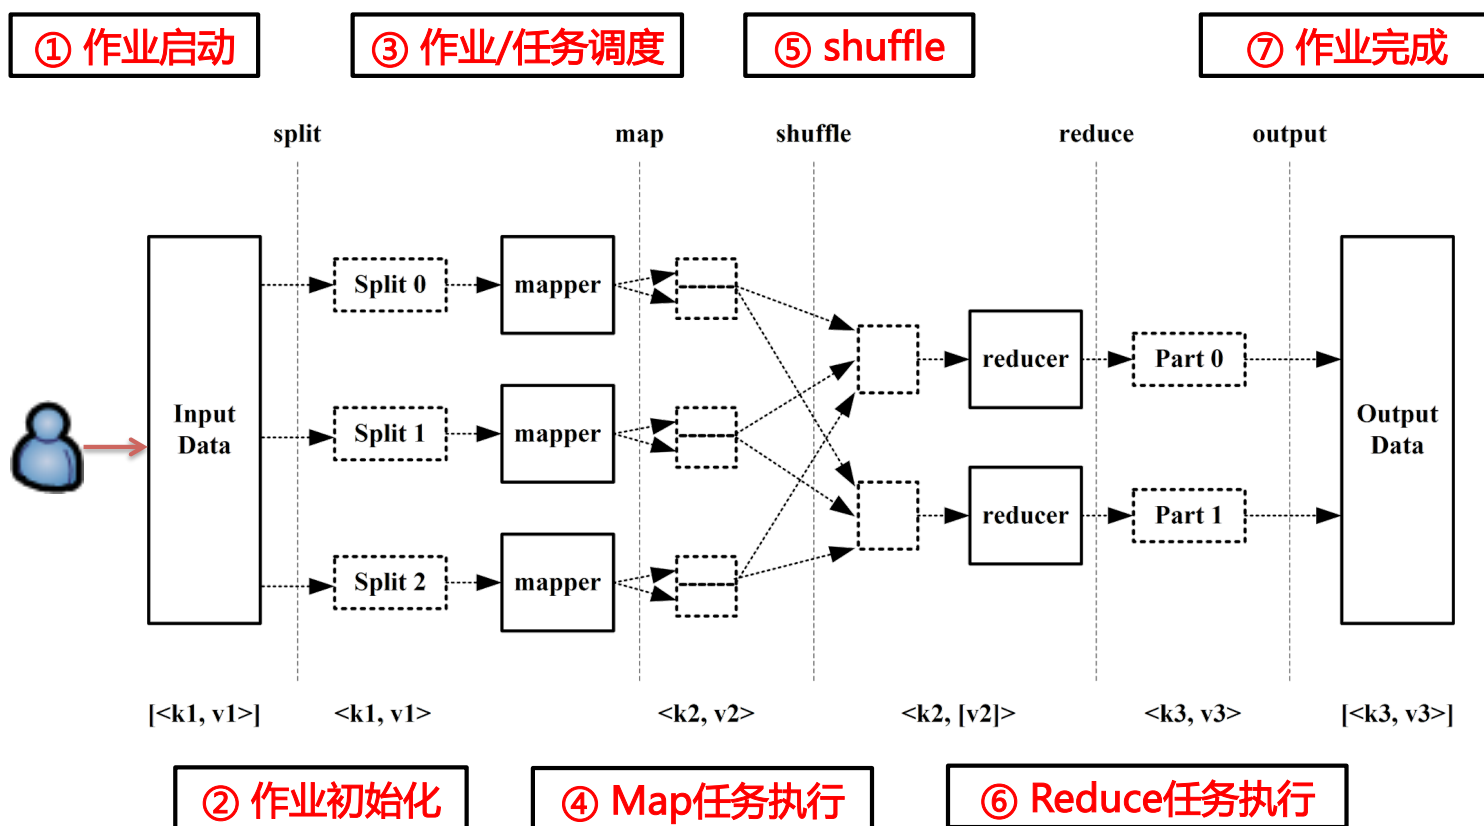


# MapReduce原理

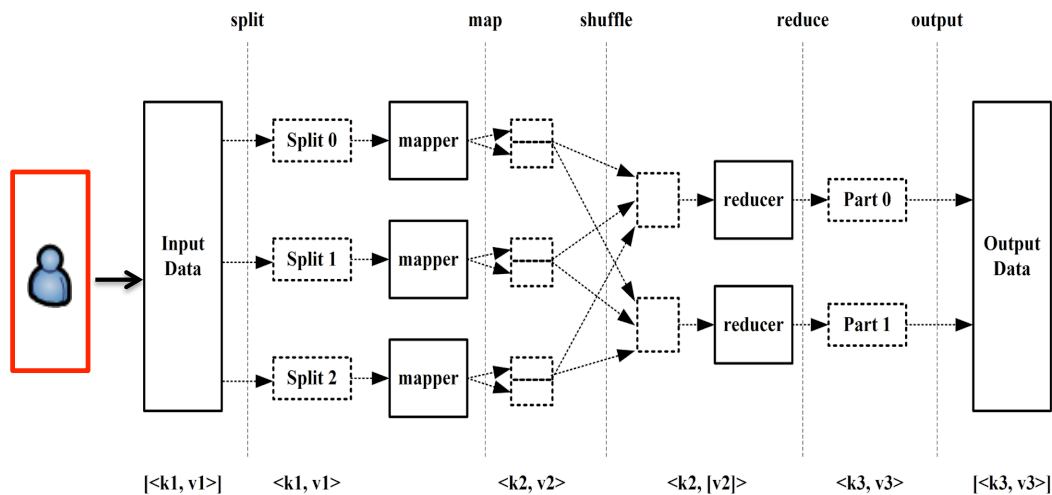
# MapReduce计算框架中的各个实体



# MapReduce作业的运行流程



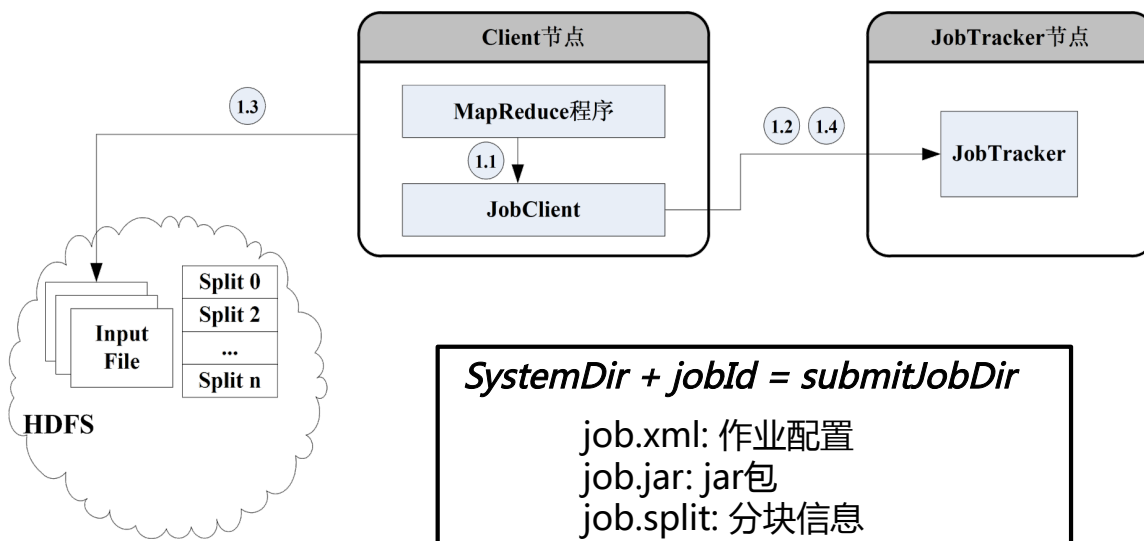
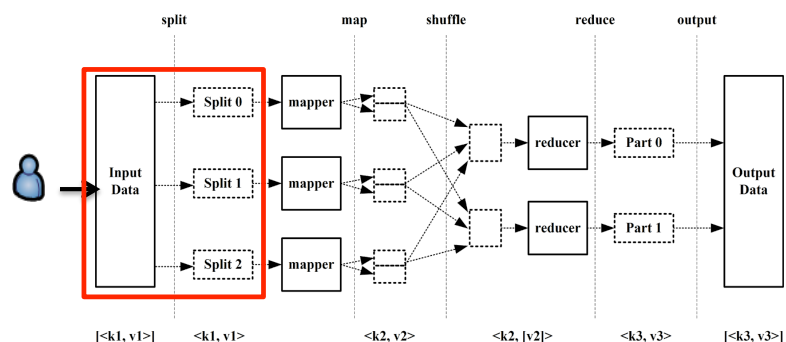
# 1. 作业启动



**`./hadoop jar ../hadoop-examples-1.0.3.jar wordcount in out`**

```
public static void main(String[] args) throws Exception{
    Configuration conf = new Configuration();
    String[] otherArgs = new GenericOptionsParser(
        conf, args).getRemainingArgs();
    if (otherArgs.length != 2) {
        System.err.println("Usage: wordcount <in> <out>");
        System.exit(2);
    }
    Job job = new Job(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job,
        new Path(otherArgs[0]));
    FileOutputFormat.setOutputPath(job,
        new Path(otherArgs[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1); }
```


## 2. 作业初始化 - 创建作业



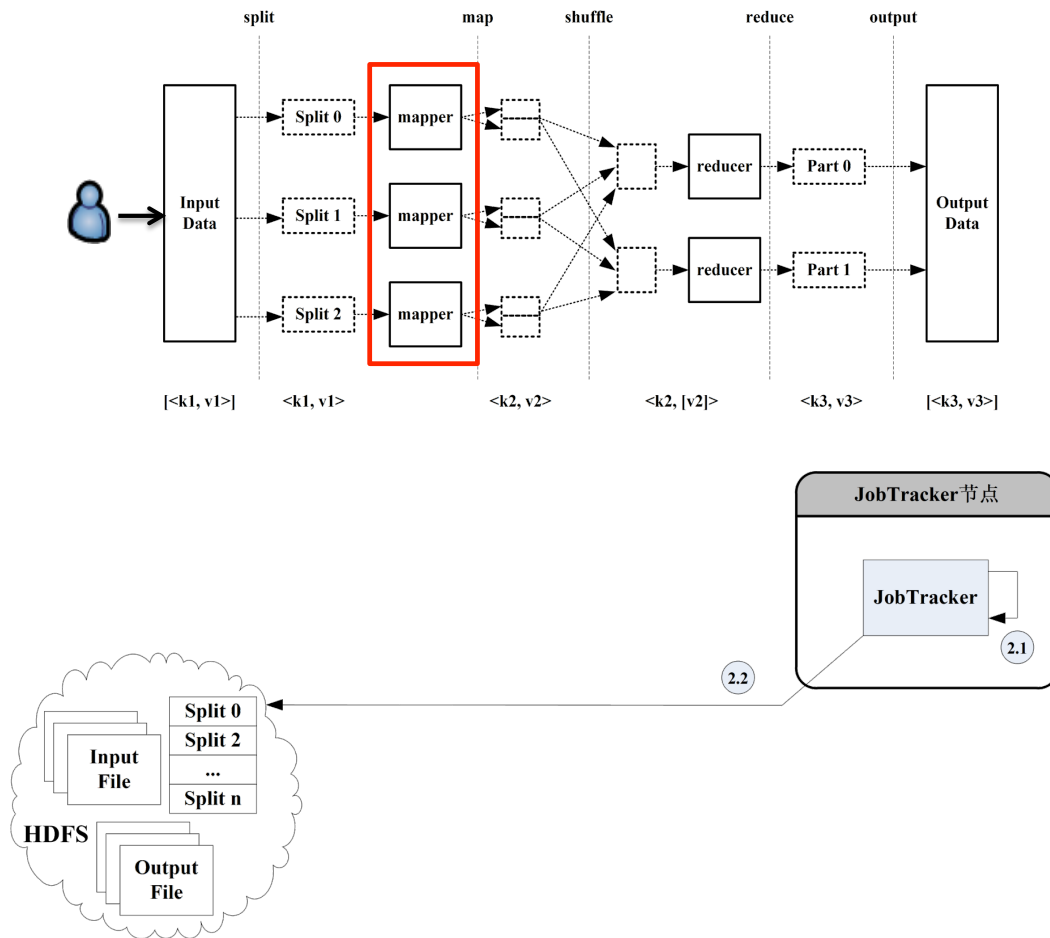
1. MapReduce 程序创建新的JobClient实例
2. JobClient向JobTracker请求获得一个新的JobId标识本次作业
3. JobClient将运行作业需要的相关资源放入作业对应的HDFS目录、计算分片数量和map任务数量
4. 向JobTracker提交作业，并获得作业的状态对象句柄



## 2. 作业初始化 - 切分数据

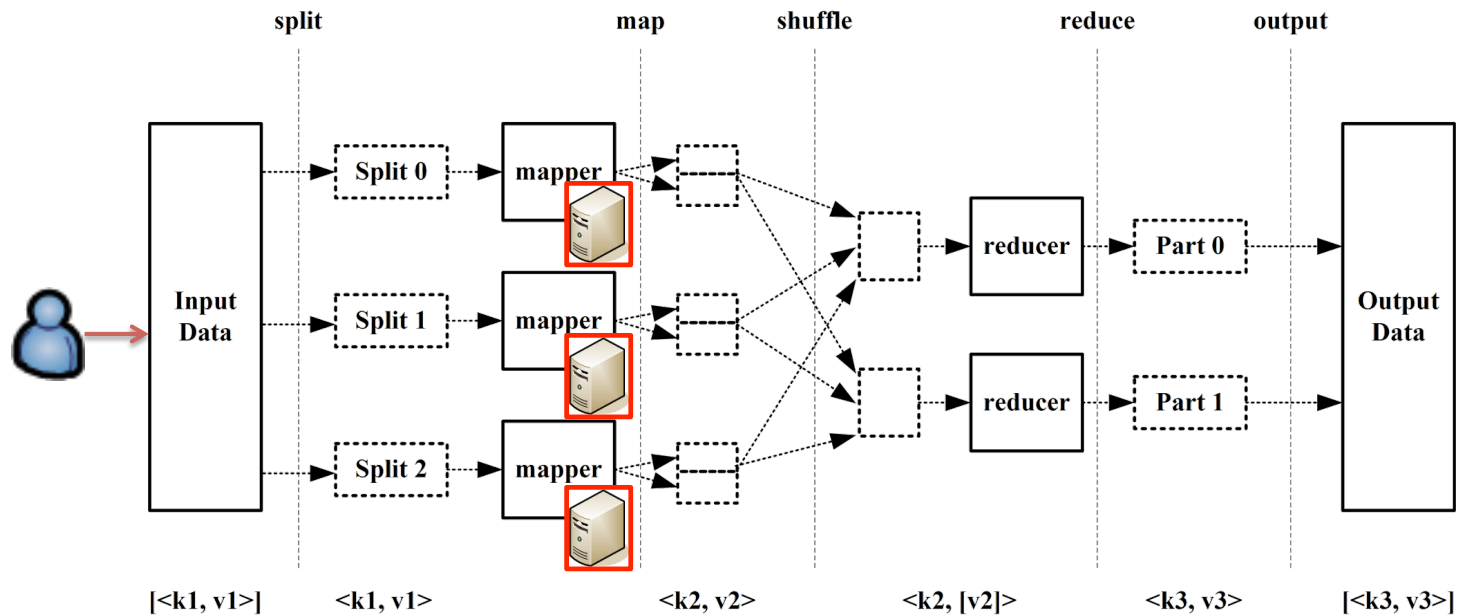


### 3. 作业 ( Job ) /任务 ( Task ) 调度

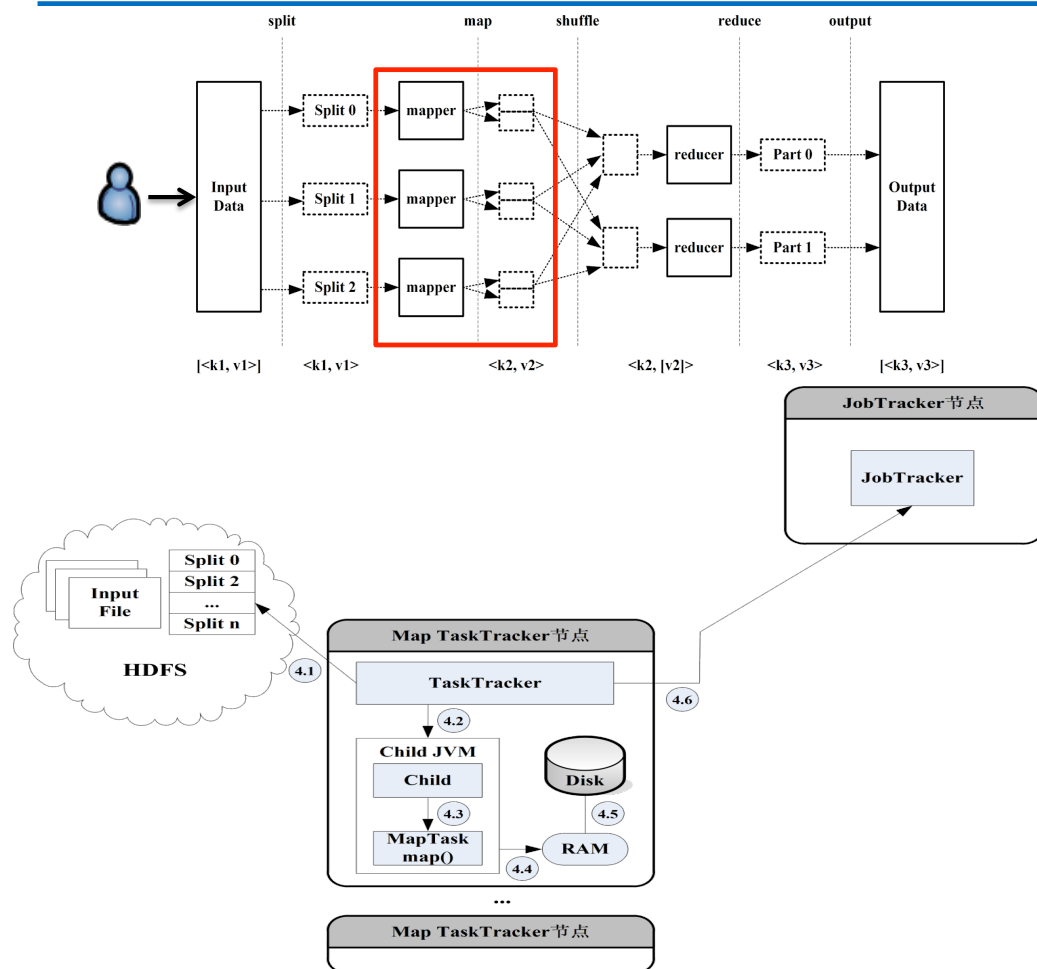


1. 作业提交请求放入队列等待调度  
， JobTracker执行调度
2. JobTracker从HDFS中取出作业分  
片信息，创建对应数量  
的TaskInProgress调度和监控Ma  
p任务

### 3. 作业 ( Job ) /任务 ( Task ) 调度完成后

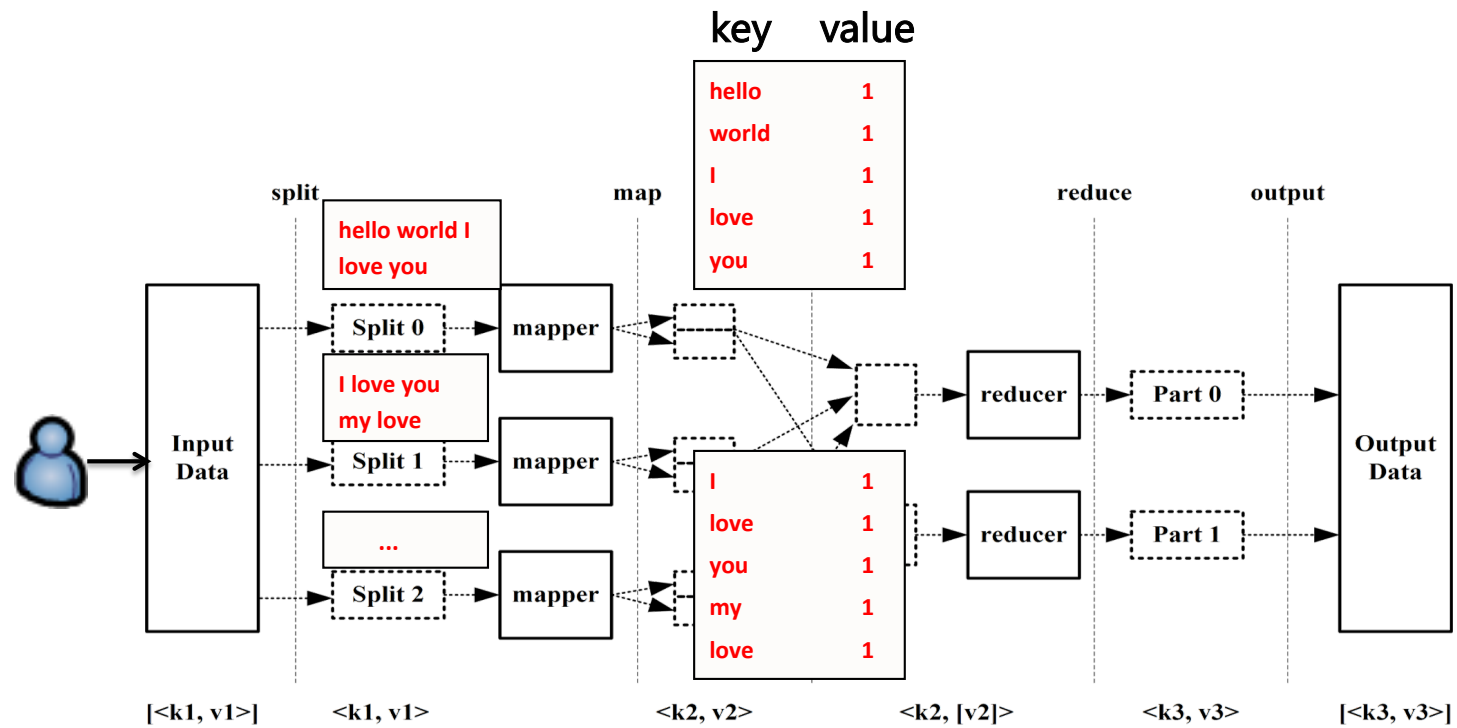


## 4. Map

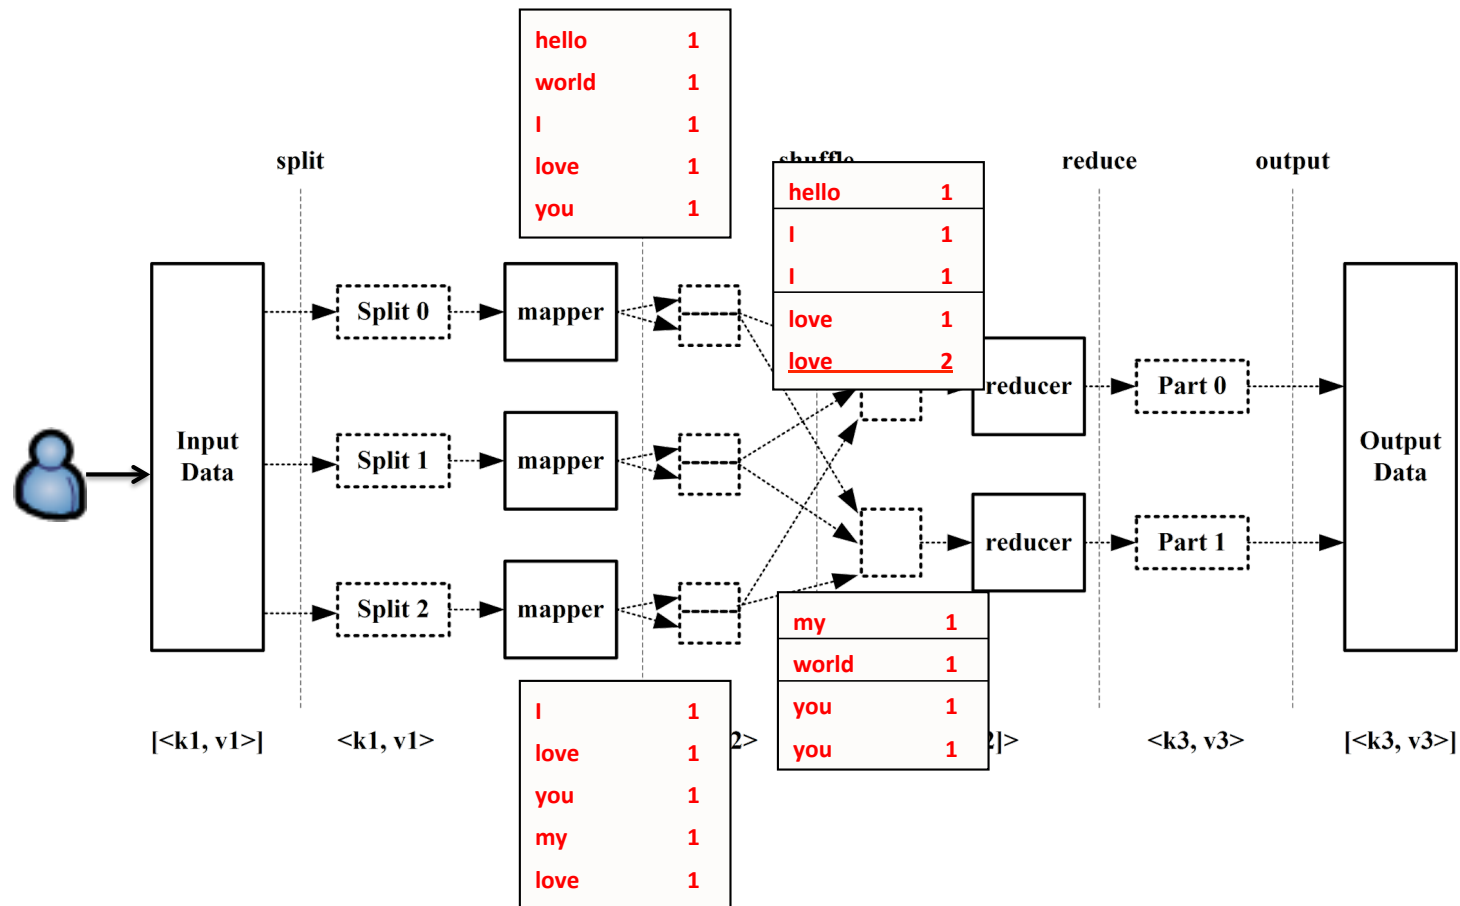


1. 从HDFS提取相关资源 ( Jar包、数据 )
2. 创建TaskRunner运行Map任务
3. 在单独的JVM中启动MapTask  
执行map函数
4. 中间结果数据定期存入缓存
5. 缓存写入磁盘
6. 定期报告进度

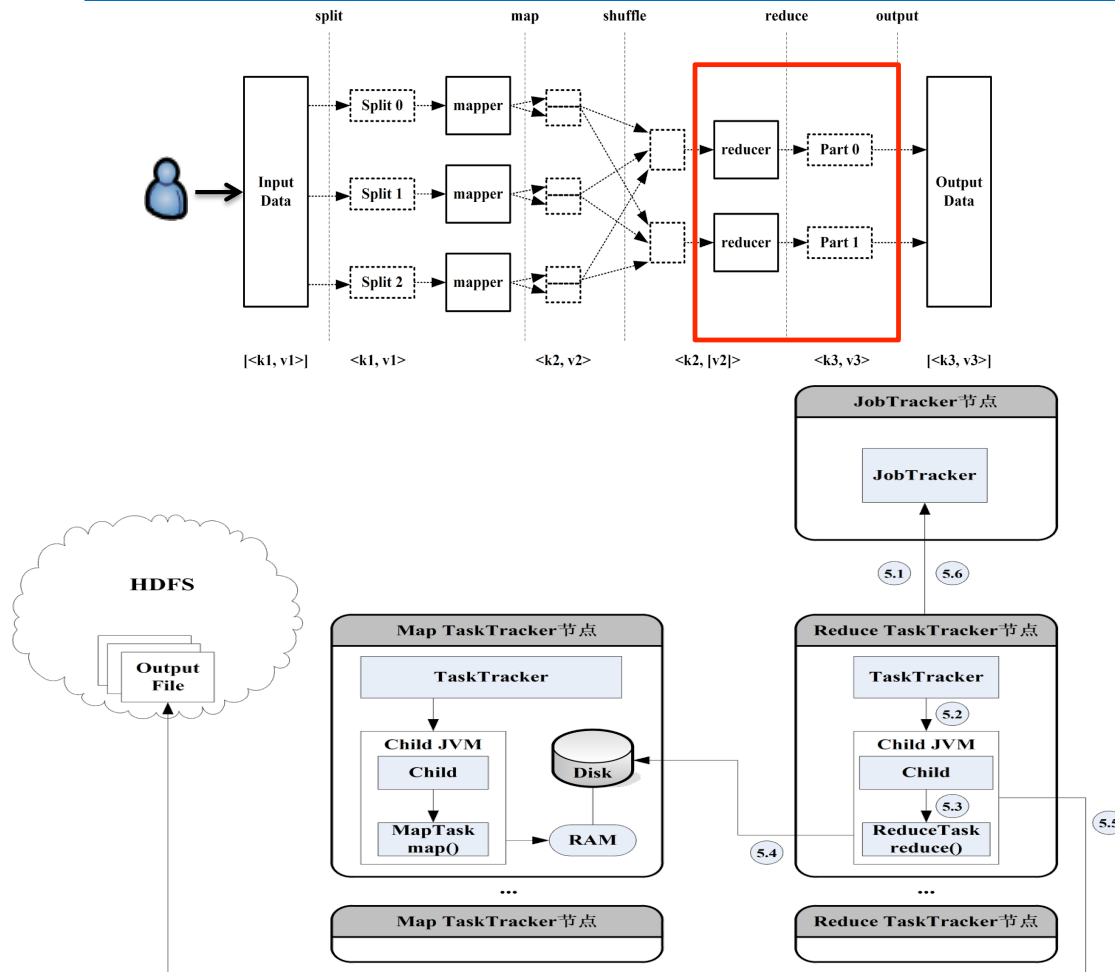
## 4. Map完成后



## 5. Shuffle

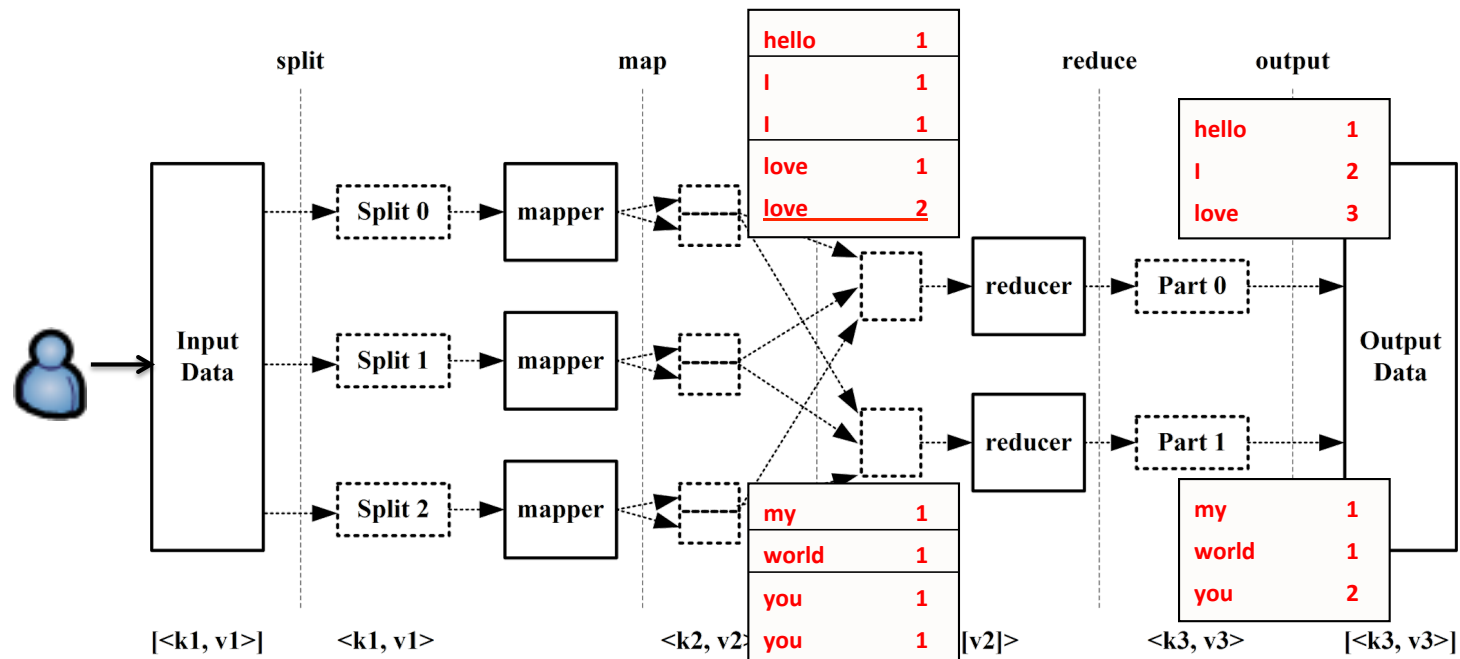


## 6. Reduce



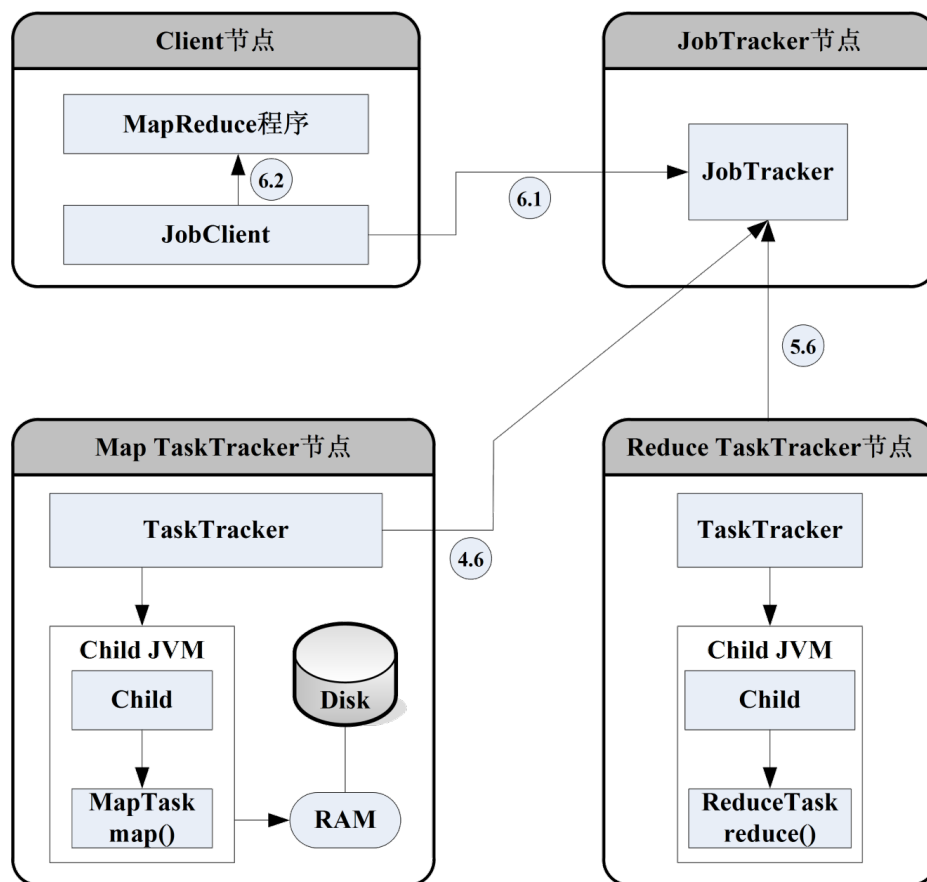
1. 分配Reduce任务
2. 创建TaskRunner运行Reduce任务
3. 在单独的JVM中启动ReduceTask  
执行reduce函数
4. 从Map节点下载中间结果数据
5. 输出结果临时文件
6. 定期报告进度

## 6. Reduce完成后





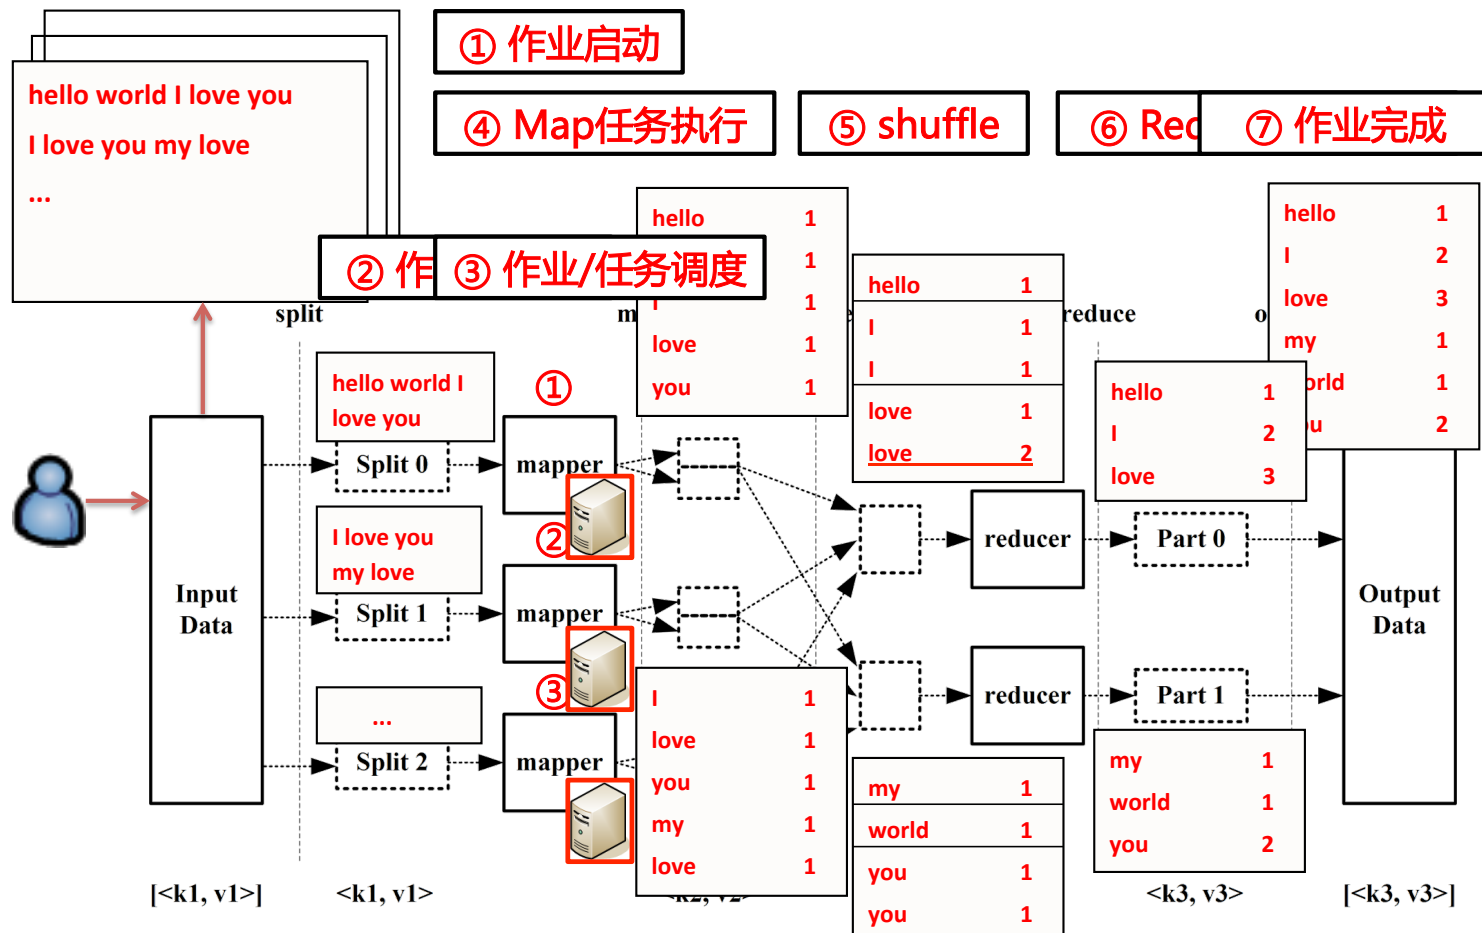
## 7. 作业完成



Map和Reduce执行时定时更新进度

1. JobClient轮询获知任务完成
2. 通知用户

# 重温MapReduce完整过程



## 下周课前请准备（实践课）

---

- 练习实践（尽力课下全部完成）：
  - HDFS上access.log文件（/home/hadooptest/Hadoop\_2016\_Autumn）
  - 编写MapReduce程序统计每个用户访问每个SP的次数、上行、下行流量
    - ✓ 在Git创建以自己名字汉语拼音全拼的分支，将代码和结果提交
    - ✓ 在wiki页面，提交运行结果图，及必要的文字说明
- （剩余工作）下周课上一起完成



刘军

北京邮电大学 信通院 宽带网络监控教研中心

北邮本部 明光楼七层

邮件地址：liujun@bupt.edu.cn

新浪微博：北邮刘军

电 话：010-62283742